

This page is for the current clarification round only. It explains Phase 1 in plain language and does not add new fixed scope unless both sides later sign a revision.

Cubecloud - Clarification Response Pack - June 2026

Pacgate AI - Phase 1 Written Clarification Reply

Thank you for the first-round questions. We have rewritten the Phase 1 explanation in plainer language. The current purchase is a local pilot and deployment-support package. It is not yet the full Pacgate commercial SaaS platform.

What This Reply Does

It focuses the real decision

- what Pacgate receives in Phase 1
- what still needs confirmation
- what should stay in later Phase 2 discussion

What This Reply Does Not Do

It does not turn every question into scope

- client portal, tenancy, billing, and platform operations stay later-stage
- external intake channels stay optional unless later added
- new commitments still require a later agreed revision

Cover Response

The cleanest reply is simple: restate the buying boundary first, then answer the details through annexes. That keeps the conversation practical and avoids a long point-by-point argument.

Commercial note: if later discussions move toward rollout-scale hardware, broader internal evaluation access, optional Feishu integration, or a seed-client charging structure, those items can be scoped as later expansion options and quoted separately if useful. They can then be reviewed and priced through a follow-on commercial revision, while the current two-machine Phase 1 purchase boundary stays as it is.

Read Phase 1 As

A practical pilot

- local deployment on agreed pilot hardware
- two reference legal-agent paths adapted for Pacgate
- starter knowledge, authority, and operating structure

Do Not Read Phase 1 As

A hidden Phase 2 commitment

- a finished external product launch
- a locked Phase 2 platform architecture
- automatic acceptance of every future platform request

Core Clarifications

Project Nature

Local pilot first

- The quoted work is a locally deployed pilot with real deliverables.
- Its job is to prove two runnable legal AI paths and support later decisions.
- It is not the full Pacgate SaaS platform build.

Delivery Shape

Concrete outputs, not abstract discovery only

- Pacgate should receive runnable environments and operating materials.
- The scope includes starter knowledge structure and bounded adaptation work.
- The output is reviewable and practical, but still pilot-grade.

Operating Model

Human review remains central

- Both system paths should be treated as human-reviewed workflows.
- Draft output, evidence check, and lawyer approval remain the default path.
- Phase 1 is meant to support lawyers, not bypass them.

Commercial Handling

Scope, advisory, and later expansion must stay separated

- Fixed delivery, advisory hours, and optional items should not be mixed together.
- Phase 2 platform options should stay separate from the current purchase.
- This keeps both buying clarity and later pricing discipline.

Specific Points That Need Correction

System 1

Do not describe Claude for Legal as a finished local product

- Treat it as a reference workflow that Cubecloud adapts for Pacgate.
- Do not imply the upstream repository will be delivered unchanged as a production-ready standalone system.

System 2

Do not describe Lavern as an already finished SaaS core

- Lavern is useful as a local-capable prototype and workflow reference.
- Phase 1 should use it as an adaptation path, not as a full production promise.

Knowledge Layer

Describe the knowledge layer as a starter structure

- Phase 1 should create Pacgate-controlled starter structure for private materials and authority tracking.
- It should not be read as a promise to finish a full legal database build or mature GraphRAG stack in the same phase.

Recommended Next Step

Step 1

Confirm the boundary wording

Align on the Phase 1 position and the split between fixed delivery, advisory support, optional add-ons, and later Phase 2 options.

Step 2

Use the annexes for detail

Use Annex A for scope, Annex B for hardware and deployment, Annex C for system and knowledge boundaries, and Annex D for later options.

Step 3

Close this round before adding new commitments

Only after the scope wording is stable should both sides decide whether any item becomes advisory work, a change request, or a separate Phase 2 track.

Supplementary Verbal Q&A

These four points are included here as supplementary reference notes for the current clarification round. They do not change the fixed Phase 1 scope unless both sides later agree a revision.

1. Commercial Expansion

Additional hardware remains a later-scale question

The current two-machine pilot is the entry configuration, not the assumed final commercial footprint. If Pacgate later moves into a broader commercial rollout, additional hardware may be needed based on user volume, concurrency, response speed, and model size. The exact model list and configuration should be fixed at contract stage. As a directional planning number only, an additional local inference workstation is currently around the low CNY 30,000s per unit including basic application deployment labor, with higher-spec options available if larger models or heavier usage are required.

2. Access And User Volume

Broader trial use should not share the tuning lane

For Phase 1, access is intentionally controlled through the service entry rather than open workstation-side embedding. Claude for Legal depends on cloud-side API services, so it should not be treated as a locally isolated workstation component in the same way as a private local model. A broader internal evaluation - for example many lawyers trying the service at once - is possible, but it should be planned as a separate access lane or separate capacity. That avoids mixing general user testing with Lavern tuning and debugging work.

3. Feishu Test Connector

An extra connector can be tested as a bounded extension

The current standard shape is one controlled service endpoint for remote use. Adding a Feishu connector for testing is technically feasible, but it would be treated as an additional integration item rather than part of the default Phase 1 baseline. The cleaner approach

is to test that connector as a bounded extension, with its own access rules and test scope.

4. Seed Client Beta

Ten-client testing starts to look like a small commercial trial

At a scale such as about 10 seed customers, the discussion is no longer only about the two-machine pilot; it becomes a small commercial trial environment. Directionally, that usually means another 2 to 3 local inference machines, each currently in the low CNY 30,000s including basic deployment labor, before choosing stronger hardware tiers. Feishu access and charging logic would add separate application and workflow work, so the final budget should be quoted after the connector flow, usage pattern, and billing method are defined.

Attachment Map

Attachment	Purpose	What it resolves
Annex A - Scope & Delivery Matrix	Shows what is fixed now, what is advisory, and what is still pending.	Prevents Phase 1 and Phase 2 from being mixed together.
Annex B - Hardware & Deployment	Explains the two-node pilot, BOM confirmation path, deployment route, and security boundary.	Pulls hardware, data-flow, and admin-access questions into one place.
Annex C - Systems & Knowledge	Explains System 1, System 2, the authority structure, and the review workflow.	Corrects the over-read of Claude for Legal, Lavern, and Package 3.
Annex D - Phase 2 Options Note	Separates future customer-entry and platform questions from the current purchase.	Keeps useful future questions alive without turning them into present commitment.

Pacgate Clarification Layer

Temporary bilingual reply surface aligned to the live proposal theme.

Use together with the annex pages and temporary workboard - not linked into the public proposal index.